

Dataset Introduction

Dataset Link: <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>

This project uses the **Telco Customer Churn** dataset from [Kaggle](#). The data was originally by **IBM** and has already been cleaned and labeled.

This project aims to:

- **perform** Exploratory Data Analysis for the dataset.
- **apply** data Preprocessing and Feature Engineering preparing the dataset for model training.
- **build a machine learning model** able to predict customer churn using their usage patterns, account information and demographics in order to help businesses take proactive retention measures.

✓ Import Dependencies

We import all necessary libraries for data manipulation, visualization and analysis.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

✓ Importing the dataset

We import the dataset from Google Drive. We simply connect our colab file to Google Drive by clicking the Drive icon on the navigation bar at the left or by running the code:

```
from google.colab import drive
drive.mount('/content/drive')
```

```
dataset = ('/content/drive/MyDrive/Machine Learning Datasets/Customer Churn Prediction/Telco-Customer-Churn.csv')
dataset = pd.read_csv(dataset)
dataset.head(5)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	Dev
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 21 columns

✓ Dataset Overview

The dataset contains information about **telecommunication company customers** and whether they **churned** (left the company) in the last month.

Target Variable: Churn (Yes = Customer left, No = customer stayed)

Column Categories

1. Target Variable

- **Churn** — Indicates if the customer left within the last month.

2. Services Signed Up

- **Phone service:** PhoneService, MultipleLines
- **Internet services:** InternetService, OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies

3. Customer Account Information

- **Tenure:** tenure — how long they've been a customer (months)
- **Contract type:** Contract
- **Billing:** PaperlessBilling
- **Payment method:** PaymentMethod
- **Charges:** MonthlyCharges, TotalCharges

4. Demographic Information

- **Gender:** gender
- **Family status:** Partner, Dependents

```
print(f"The dataset has {dataset.shape[0]:,} rows and {dataset.shape[1]:,} columns.")
dataset.head(5)
```

The dataset has 7,043 rows and 21 columns.

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	Dev
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 21 columns

▼ Data Types and Missing Values

Inspecting column data types and identifying missing values for cleaning and preprocessing.

dataset.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
Column Non-Null Count Dtype
--- -
0 customerID 7043 non-null object
1 gender 7043 non-null object
2 SeniorCitizen 7043 non-null int64
3 Partner 7043 non-null object
4 Dependents 7043 non-null object
5 tenure 7043 non-null int64
6 PhoneService 7043 non-null object
7 MultipleLines 7043 non-null object
8 InternetService 7043 non-null object
9 OnlineSecurity 7043 non-null object
10 OnlineBackup 7043 non-null object
11 DeviceProtection 7043 non-null object
12 TechSupport 7043 non-null object
13 StreamingTV 7043 non-null object
14 StreamingMovies 7043 non-null object
15 Contract 7043 non-null object
16 PaperlessBilling 7043 non-null object
17 PaymentMethod 7043 non-null object
18 MonthlyCharges 7043 non-null float64

```

19 TotalCharges      7043 non-null  object
20 Churn              7043 non-null  object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
dataset.isnull().sum()
```



	0
customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0

```
dtype: int64
```


▼ Class Distribution

We check for class imbalance. If class count differ significantly a possible use case scenario for SMOTE might be needed later on to normalize the class distribution.

```

dataset['Churn'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()

```



Statistical Summary

We use the `.describe()` method to obtain a statistical summary of the numerical features. This includes measures such as mean, standard deviation, and quartiles, which help assess the central tendency and spread of data.

```
dataset.describe()
```



	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

Unique Values for Categorical Features

Checking different categorical values in each columns. It helps in planning encoding strategies later on in the process.

```
cat_cols = dataset.select_dtypes(include=['object']).columns
```

```
for col in cat_cols:
    print(f"{col}: {dataset[col].unique()}\n")
```



```
customerID: ['7590-VHVEG' '5575-GNVDE' '3668-QPYBK' ... '4801-JJAZL' '8361-LTMKD'
'3186-AJIEK']

gender: ['Female' 'Male']

Partner: ['Yes' 'No']

Dependents: ['No' 'Yes']

PhoneService: ['No' 'Yes']

MultipleLines: ['No phone service' 'No' 'Yes']

InternetService: ['DSL' 'Fiber optic' 'No']

OnlineSecurity: ['No' 'Yes' 'No internet service']
```

```

OnlineBackup: ['Yes' 'No' 'No internet service']
DeviceProtection: ['No' 'Yes' 'No internet service']
TechSupport: ['No' 'Yes' 'No internet service']
StreamingTV: ['No' 'Yes' 'No internet service']
StreamingMovies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
TotalCharges: ['29.85' '1889.5' '108.15' ... '346.45' '306.6' '6844.5']
Churn: ['No' 'Yes']

```

Observation: We can observe that on the given list of Categorical columns `TotalCharges` was included. This could mean that the values were stored as an object string instead of a numeric type. In order to solve this we simply convert the whole column into numeric type.

✓ Converting `TotalCharges` into numeric type

We simply convert the whole column by using `pd.to_numeric`.

```

print("Initial column type:", dataset['TotalCharges'].dtypes)
dataset['TotalCharges'] = pd.to_numeric(dataset['TotalCharges'], errors='coerce')
print("Total NaN values:", dataset['TotalCharges'].isna().sum())

```

```

↗ Initial column type: object
Total NaN values: 11

```

```

dataset['TotalCharges'] = dataset['TotalCharges'].fillna(0)
print("Final column type:", dataset['TotalCharges'].dtypes)
print("Total NaN values:", dataset['TotalCharges'].isna().sum())

```

```

↗ Final column type: float64
Total NaN values: 0

```

✓ Exploratory Data Analysis

✓ Numerical Feature Distribution

Visualizing numerical features such as `tenure`, `MonthlyCharges`, and `TotalCharges` helps us understand:

- the range and spread of each variable.
- presence of outliers.
- General customer spending and tenure patterns.

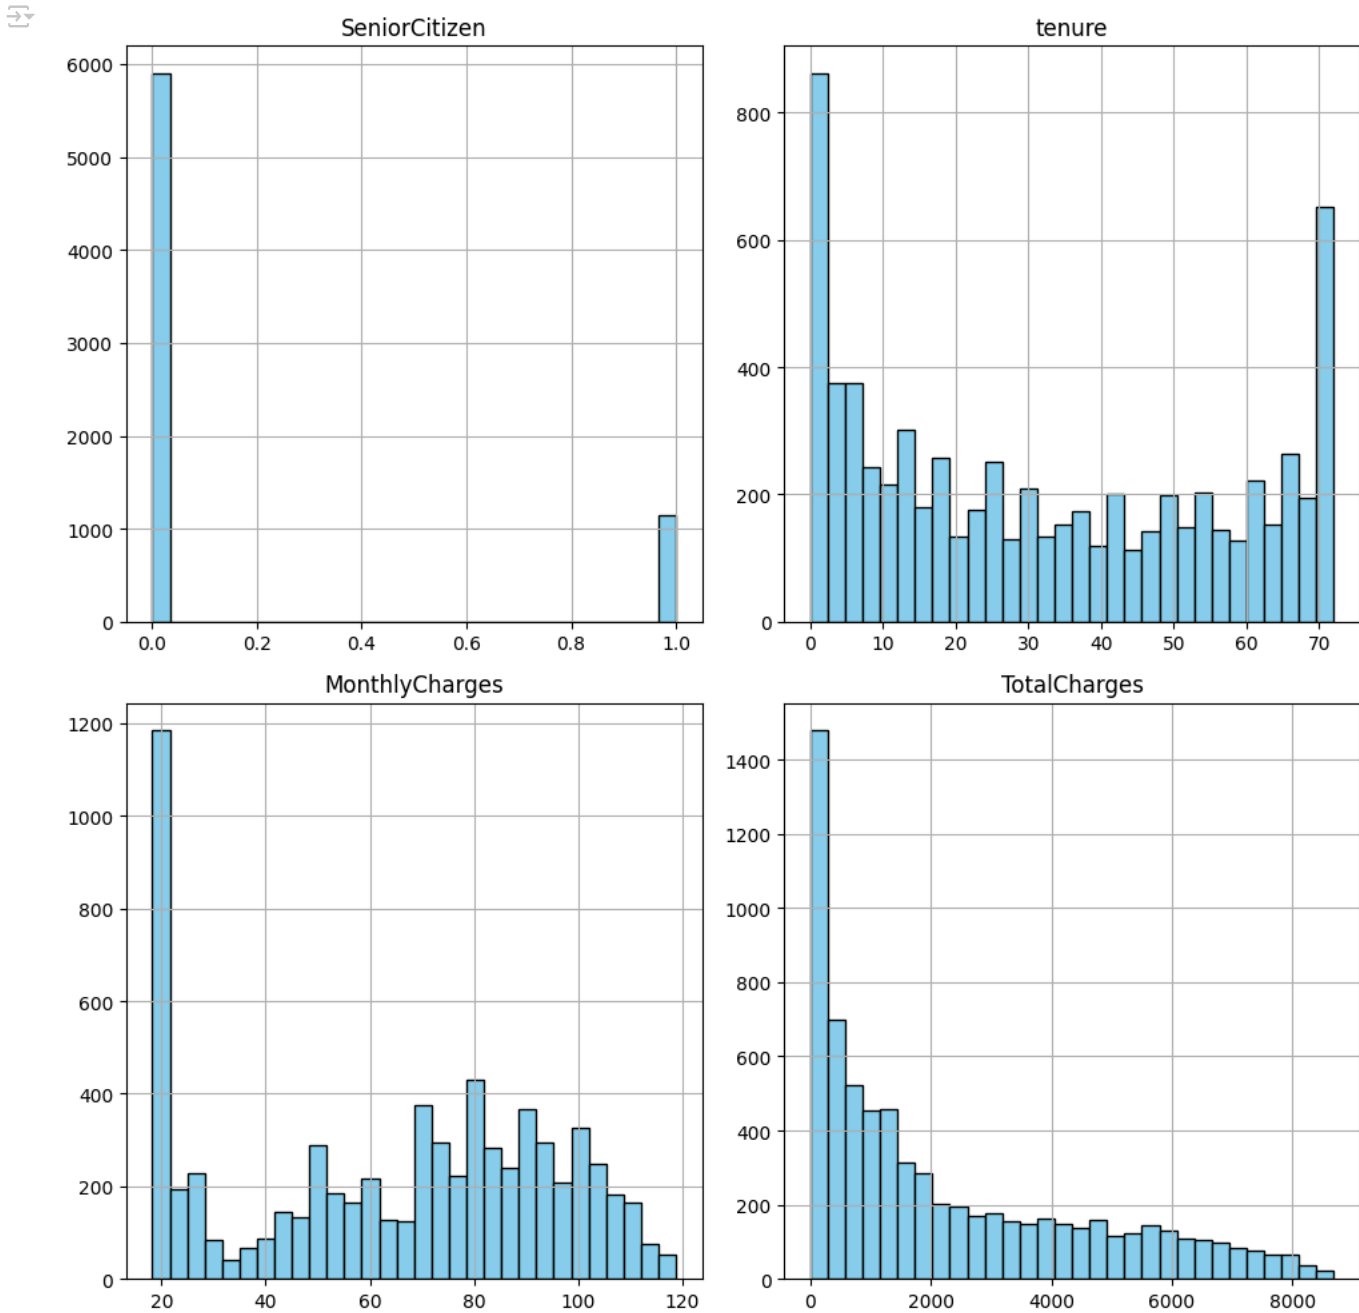
This step is useful for spotting potential skewness and deciding whether transformations (e.g., log-scaling) might be needed.

```

num_cols = dataset.select_dtypes(include=['int64', 'float64']).columns

dataset[num_cols].hist(bins=30, figsize=(10, 10), color='skyblue', edgecolor='black')
plt.tight_layout()
plt.show()

```



▼ Categorical Feature Distributions

Examine the distribution of categorical features such as `Contract`, `InternetService`, and `PaymentMethod`. This will show us:

- the most common customer service type
- whether any categories are rare and might need combining or special handling.

```
cat_cols = dataset.select_dtypes(include=[object]).columns.drop(['Churn', 'customerID'])

fig, axes = plt.subplots(nrows=(len(cat_cols) + 2) // 3, ncols=3, figsize=(15, 15))
axes = axes.flatten()

for i, col in enumerate(cat_cols):
    sns.countplot(y=col, data=dataset, ax=axes[i], palette='pastel', hue=col, legend=False)
    axes[i].set_title(f'{col} Distribution')
    axes[i].set_ylabel('')
    axes[i].set_xlabel('')

for j in range(len(cat_cols), len(axes)):
    fig.delaxes(axes[j])
```

```
plt.tight_layout()  
plt.show()
```




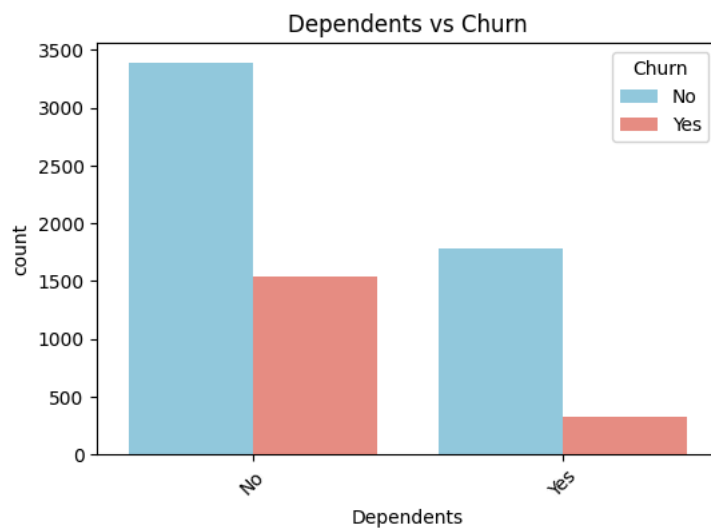
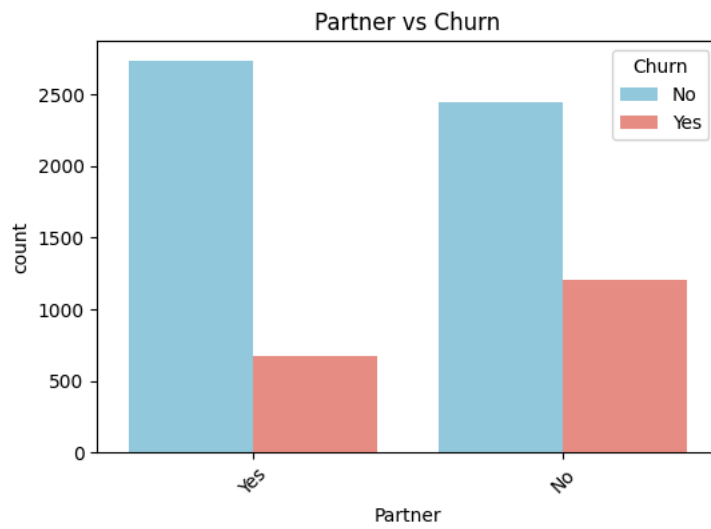

✓ Target Variable vs Feature correlation

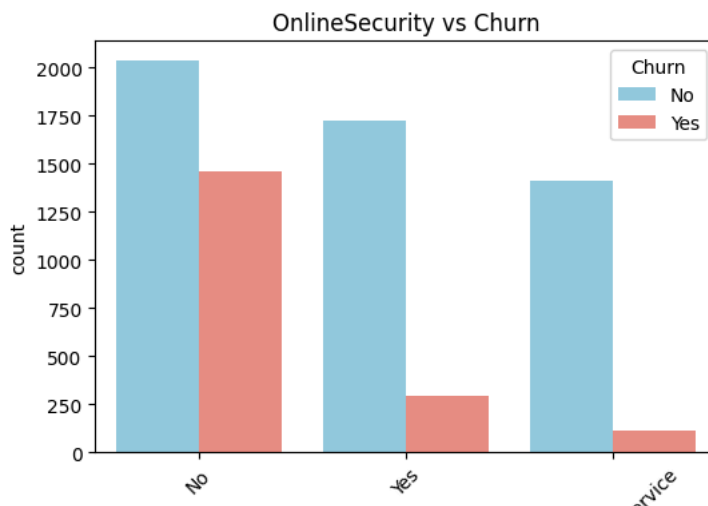

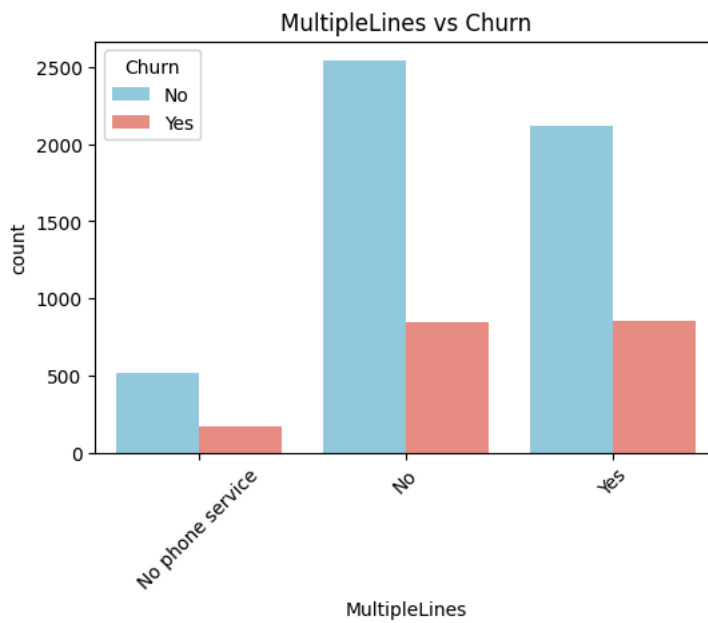

We now explore how each feature relates to churn.

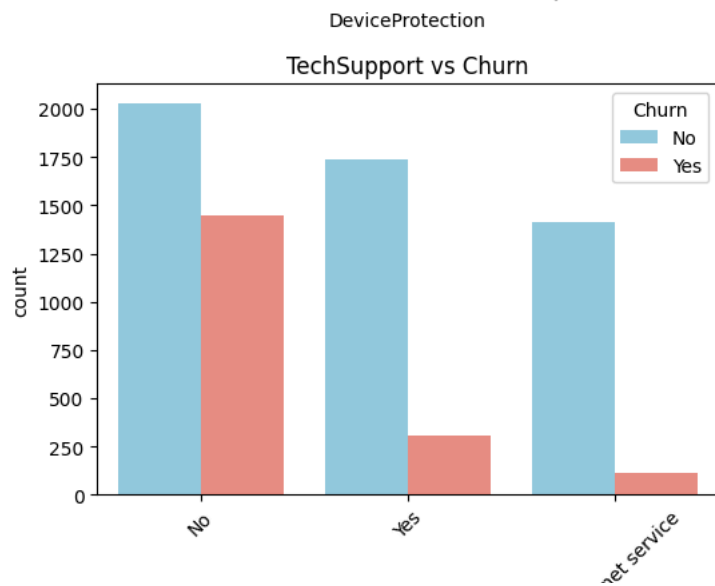
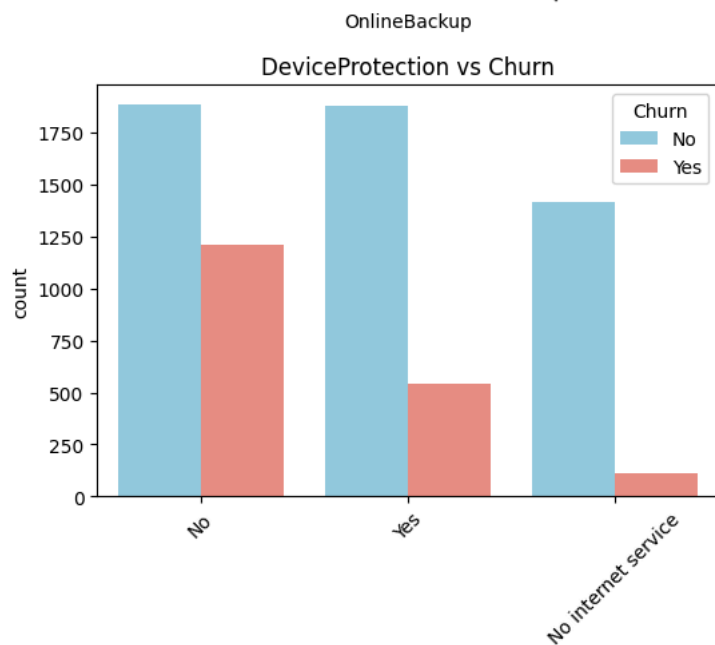
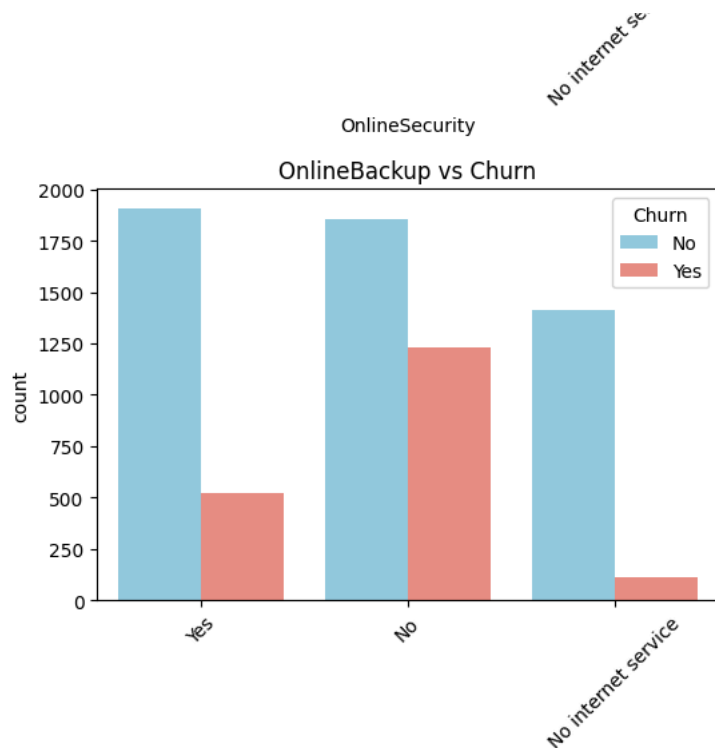
- For categorical variables: Compare category frequencies between churned and non-churned customers.
- For numerical variables: Compare value distributions using boxplots.

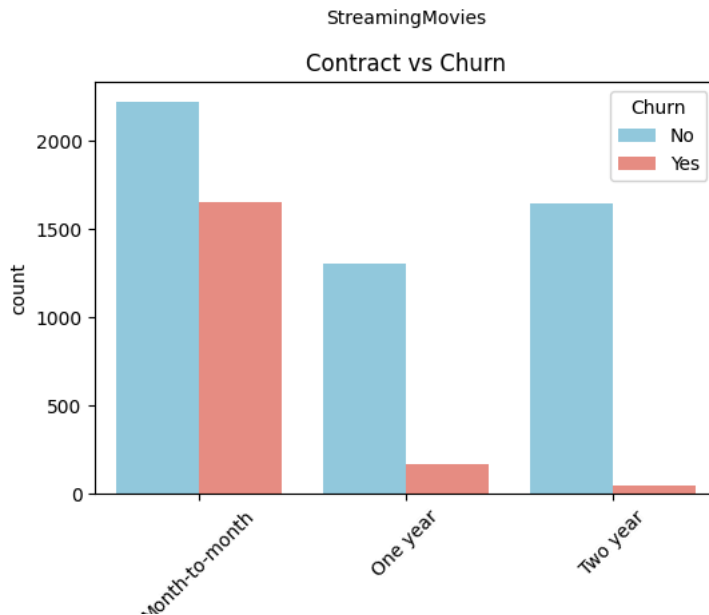
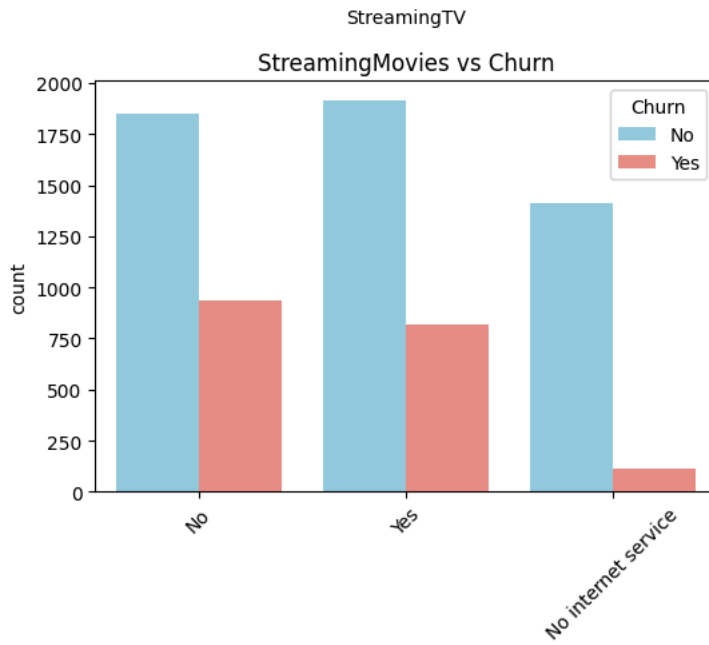

This helps identify key churn drivers, such as whether month-to-month contracts or high monthly charges are linked to higher churn rates.



```
for col in cat_cols:
    plt.figure(figsize=(6,4))
    sns.countplot(x=col, hue='Churn', data=dataset, palette=['skyblue', 'salmon'])
    plt.title(f'{col} vs Churn')
    plt.xticks(rotation=45)
    plt.show()
```












Correlation Analysis

We now check the relationship between numerical variable using correlation heatmap. This helps:

- Identify redundant features (high correlation, rule of thumb: $|\text{correlation}| > 0.85$ consider dropping one feature).
- Spot variables with potential predictive power for churn.

For example, MonthlyCharges and TotalCharges might be strongly correlated, meaning we should be cautious about multicollinearity.

```
plt.figure(figsize=(8,6))
corr = dataset[num_cols].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



Data Processing

We now proceed to **Data Processing** ensuring the data is clean, numeric where necessary and ready for ML models.

```
dataset = dataset.drop('customerID', axis=1)
dataset.head()
```



	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Device
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No	
4	Female	0	No	No	2	Yes	No	Fiber optic	No	No	

Encoding Categorical Variables

Machine learning models require numerical input, so categorical features must be transformed into numeric format. In this dataset, we have three types of categorical features:

- 1. **Binary (Yes/No)** — Converted directly to 0 and 1.
- 2. **Ordinal (ordered)** — Categories have a natural order, encoded using **Ordinal Encoding**.
- 3. **Nominal (unordered)** — No inherent order, encoded using **One-Hot Encoding** to create separate binary columns for each category.

This ensures the data is in a suitable format for all models, avoids misinterpretation of category order, and preserves model flexibility.

Encoding Binary Yes/No Columns

```
binary_map = {'Yes': 1, 'No': 0}
binary_cols = ['Partner', 'Dependents', 'PhoneService', 'PaperlessBilling', 'Churn']

for col in binary_cols:
    dataset[col] = dataset[col].map(binary_map)

dataset['gender'] = dataset['gender'].map({'Male': 1, 'Female': 0})
```

Encoding Ordinal Columns

```
from sklearn.preprocessing import OrdinalEncoder


ordinal_cols = ['Contract']
ordinal_order = [['Month-to-month', 'One year', 'Two year']]
dataset[ordinal_cols] = OrdinalEncoder(categories=ordinal_order).fit_transform(dataset[ordinal_cols])
```

One-hot encoding Nominal Columns

```
nominal_cols = ['InternetService', 'PaymentMethod', 'MultipleLines',
                'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
                'TechSupport', 'StreamingTV', 'StreamingMovies']

dataset = pd.get_dummies(dataset, columns=nominal_cols, drop_first=True)

dataset.head()
```



	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Contract	PaperlessBilling	MonthlyCharges	TotalCharges	...	Online
0	0	0	1	0	1	0	0.0	1	29.85	29.85	...	
1	1	0	0	0	34	1	1.0	0	56.95	1889.50	...	
2	1	0	0	0	2	1	0.0	1	53.85	108.15	...	
3	1	0	0	0	45	0	1.0	0	42.30	1840.75	...	
4	0	0	0	0	2	1	0.0	1	70.70	151.65	...	

5 rows × 30 columns

✓ Checking Numerical Feature Distributions

Before proceeding and deciding on a scaling method, it's important to check the distribution of numerical features. This helps us see:

- If features are **roughly normally distributed** → `StandardScaler` is a good choice.
- If features are **skewed** → we may need `MinMaxScaler` or transformations (e.g., log).
- If features have **extreme outliers** → `RobustScaler` may be better.

We'll plot histograms for each numeric feature to visualize their distributions.

```
numeric_features = dataset.select_dtypes(include=['int64', 'float64']).columns.drop(['Churn'])

plt.figure(figsize=(15, len(numeric_features) * 3))
```



```
for i, col in enumerate(numeric_features, 1):
    plt.subplot(len(numeric_features), 1, i)
    sns.histplot(dataset[col], kde=True, bins=30)
    plt.title(f'{col} Distribution')
    plt.xlabel(col)
    plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```

